

Catherine Hawley

4/7/2025

Project Two Conference Presentation: Cloud Development

https://youtu.be/wXA_gtUAJBk

Slide 1

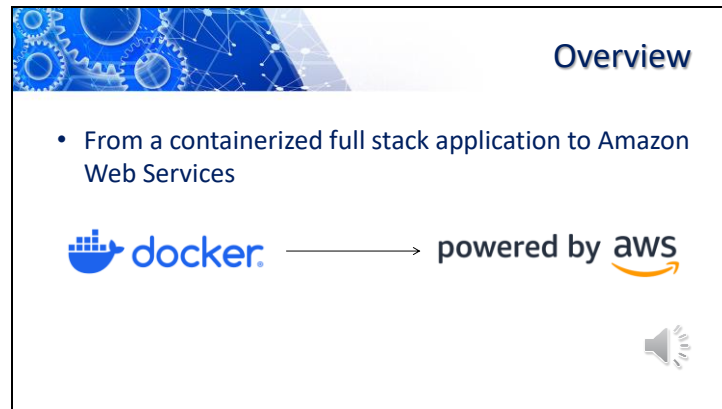


CS 470 Project Two
Conference Presentation:
Cloud Development

Catherine Hawley
April 2014



My name is Catherine Hawley, and I'm a few short months away from earning my Bachelor's degree in Computer Science. This is my presentation of my final project for CS 470, Full Stack Development Two.


Slide 2



Overview

- From a containerized full stack application to Amazon Web Services

 docker. —————> powered by  aws

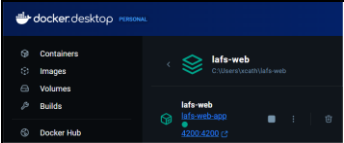


Over the last few years I've done a lot of development work for local programs, from quick functions to entire applications. In the last few weeks, I've taken an existing full stack site and migrated it to a cloud-native web application using AWS microservices.

Images: Docker Brand Guidelines, <https://www.docker.com/company/newsroom/media-resources/>

AWS Trademark Guidelines, <https://aws.amazon.com/trademark-guidelines/>


Slide 3



Containerization


Models for Migrating

1. Lift and shift (Rehost)
2. Refactor (Cloud-Native)



Tools for Containerization

- Docker Desktop, Compose
- Visual Studio Code (IDE)
- MEAN Stack
 - MongoDB
 - Express
 - Angular
 - NodeJS



To move an application to the cloud, there are two options. Lift and shift is the non-invasive way to move the pieces of the application onto the cloud, changing the application as little as possible. Refactoring involves rethinking the application, changing its structure or functions to better fit and take advantage of the cloud capabilities. For this foray into cloud-based applications, we used the lift and shift method.

This application began with Docker, a container-based virtualization tool. Docker Compose was used to control those containers. The coding was done with Visual Studio Code, and utilized the MEAN stack of MongoDB, Express, Angular, and NodeJS.

Images: Screenshot of my instance of Docker Desktop

Strategies for Lift & Shift Cloud Migration, <https://www.linkedin.com/pulse/strategies-lift-shift-cloud-migration-irene>


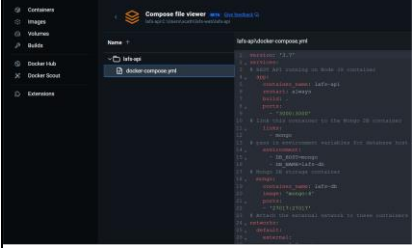
MEAN Stack Best Practices, <https://medium.com/@sureshkrishna75525/mean-stack-best-practices-pro-tips-for-efficient-development-in-2023-e035168eff25>

Slide 4

Orchestration

The value of using Docker Compose

- Run containers simultaneously
- Program scalability
- Apply settings to all containers
 - Keeps containers consistent
 - Allows containers to communicate



Docker Compose allows you to treat all of your containers as a single program, which is helpful because it is. The database container run by Mongo and the API container running on Node can be started and stopped simultaneously using Compose. Compose also offers scalability for container-based applications, which would be complicated if many containers needed to be managed individually.

Settings can be applied to all containers from the docker-compose YAML file, a single convenient place. This ensures settings are compatible between containers, and gives us networking capability, allowing the containers to communicate.

The slide is titled "The Serverless Cloud" in a blue font at the top right. The top left corner features a blue graphic with white gears and network lines. The main content area is white and contains a section titled "Serverless" in bold. Below this title are two bullet points. The first bullet point is "What is 'serverless'?" followed by two sub-points: "More like 'some one else's servers'" and "Cost-saving, time-saving, auto-scaling, and division of labor." The second bullet point is "S3 = Simple Storage Service" followed by four sub-points: "Serverless storage solution", "AWS integration", "Intelligent Tiering", and "Backup redundancies". A small speaker icon is located in the bottom right corner of the slide.

The Serverless Cloud

Serverless

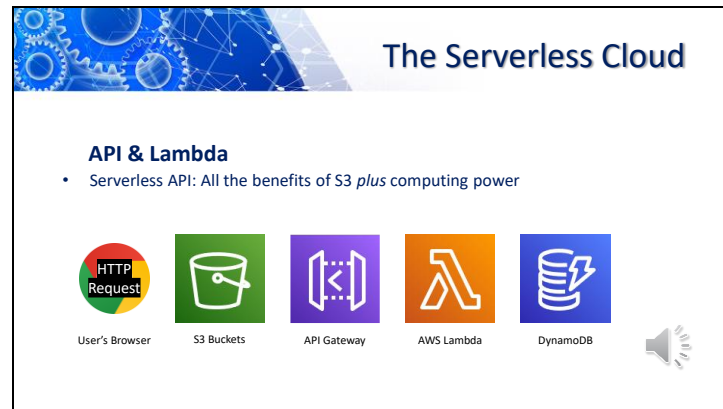
- What is "serverless"?
 - More like "some one else's servers"
 - Cost-saving, time-saving, auto-scaling, and division of labor.
- S3 = Simple Storage Service
 - Serverless storage solution
 - AWS integration
 - Intelligent Tiering
 - Backup redundancies
 - Versioning

What IS serverless?

By "serverless" we mean that our program runs on someone else's servers. We're using Amazon Web Services as our cloud provider, so Amazon owns the hardware that we're using. Server hardware is expensive in terms of time and money, as it requires an initial purchase, setup, and maintenance, along with decisions regarding scale and structure. Allowing AWS to handle those costs and decisions can save a business or individual developer money and time, allowing them to focus on developing the product.

S3 means Simple Storage Service, Amazon's cloud based storage solution. It brings in all the benefits of serverless architecture, links to other AWS technologies, and a few other S3 perks. Intelligent tiering is a cost-saving measure for less frequently accessed data, backups in multiple locations offer redundancy in case of physical server damage, and versioning allows access to previous versions of the same file.

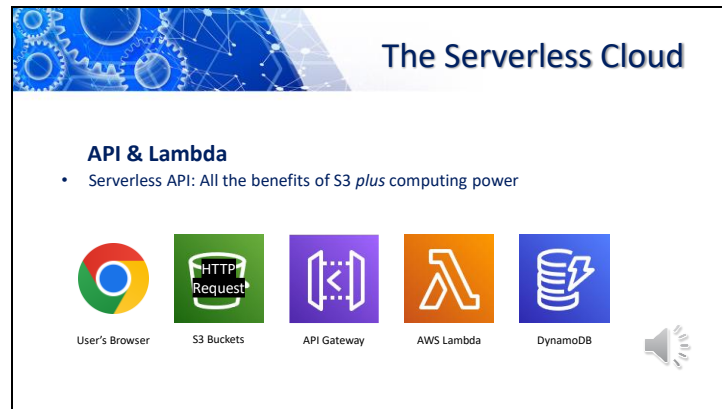
Slide 6



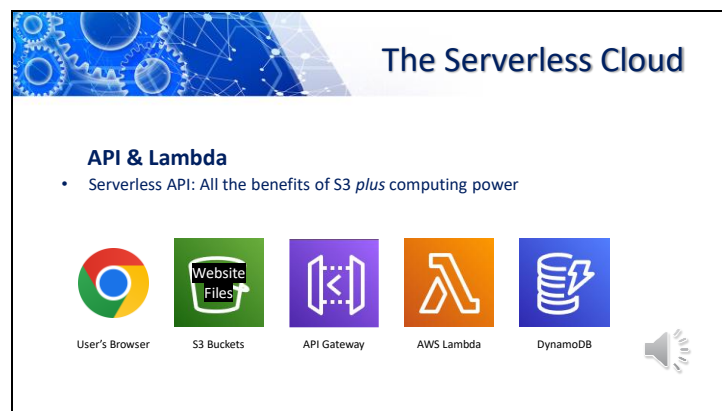
Lambda is event-driven computing. We use NodeJS to write the scripts for our Lambda functions. Each function gets its own block of code, separate from the other functions. This standalone method keeps the code maintainable and understandable, and ensures that each function can be correctly triggered by the API Gateway. The frontend and the backend are integrated by Cross-Origin Resource Sharing, or CORS, which allows our website to pull information through the API in a secure way.

The user's browser sends an HTTP request to the S3 bucket, which returns website files. Then, the user is able to interact with the website. That interaction is sent to the API Gateway in the form of an HTTP request. The gateway routes that request as an event for Lambda. Our Lambda functions are all CRUD operations, which are sent to the database. Dynamo processes that operation and sends a response back through Lambda and the API gateway to the user's browser, giving them the requested information or letting them know the database has been updated.

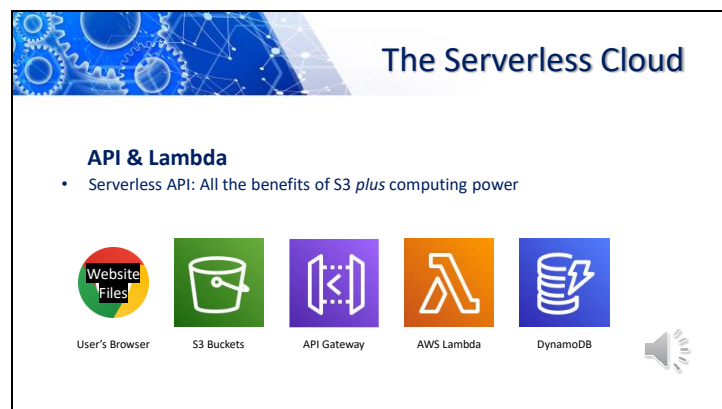
Slide 7



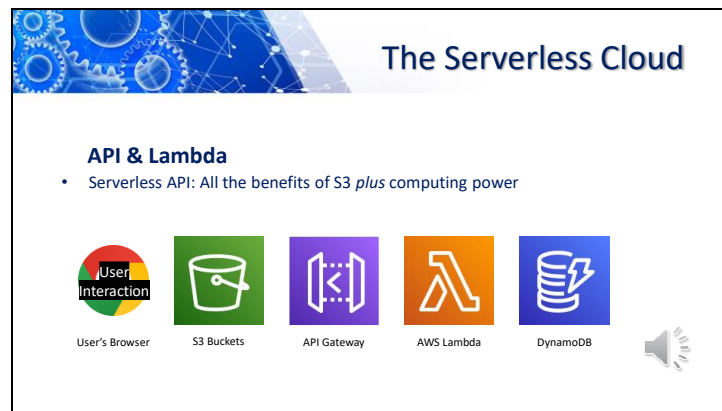
Slide 8



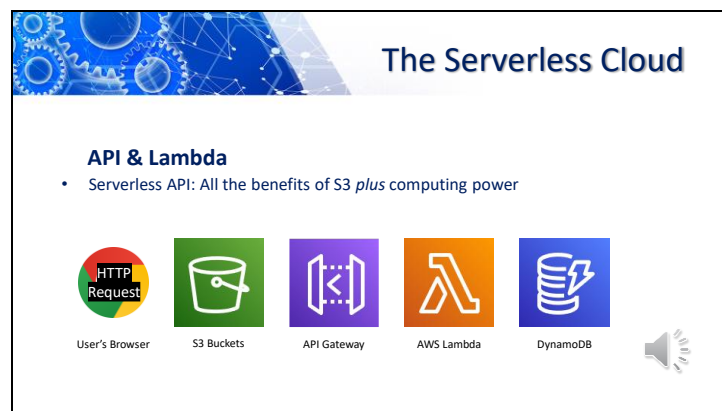
Slide 9



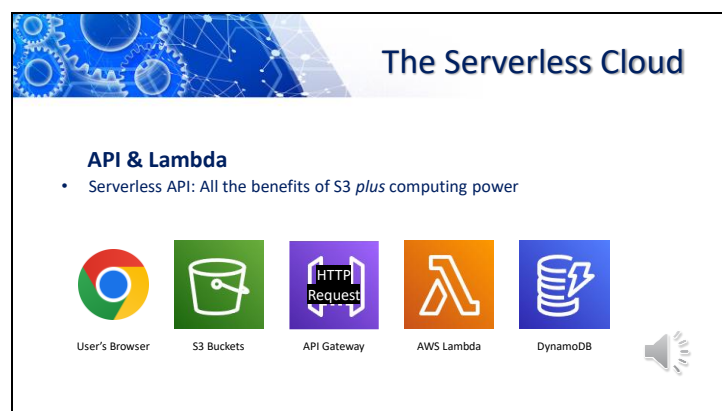
Slide 10



Slide 11



Slide 12



Slide 13

The Serverless Cloud

API & Lambda

- Serverless API: All the benefits of S3 *plus* computing power

The diagram illustrates a serverless architecture flow. It starts with a 'User's Browser' (Google Chrome icon), which triggers an event in 'S3 Buckets' (bucket icon). This event is then processed by 'API Gateway' (gate icon with 'Event' label), which calls 'AWS Lambda' (lambda icon). The lambda function interacts with 'DynamoDB' (cylinder icon with lightning bolt). A speaker icon is present at the bottom right.

Slide 14

The Serverless Cloud

API & Lambda

- Serverless API: All the benefits of S3 *plus* computing power

This slide is identical to Slide 13, showing the same serverless architecture flow: User's Browser -> S3 Buckets -> API Gateway -> AWS Lambda -> DynamoDB. A speaker icon is present at the bottom right.

Slide 15

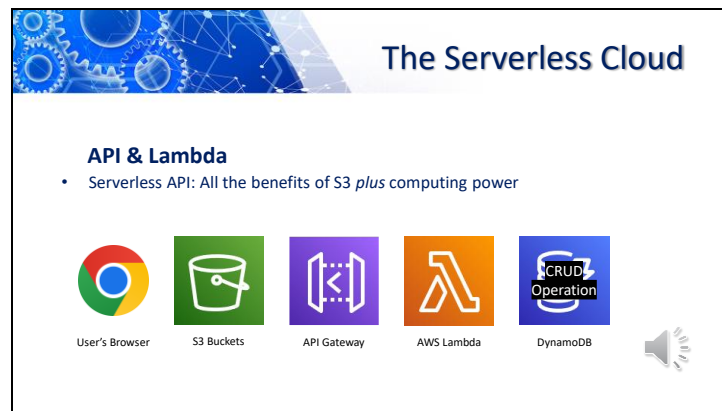
The Serverless Cloud

API & Lambda

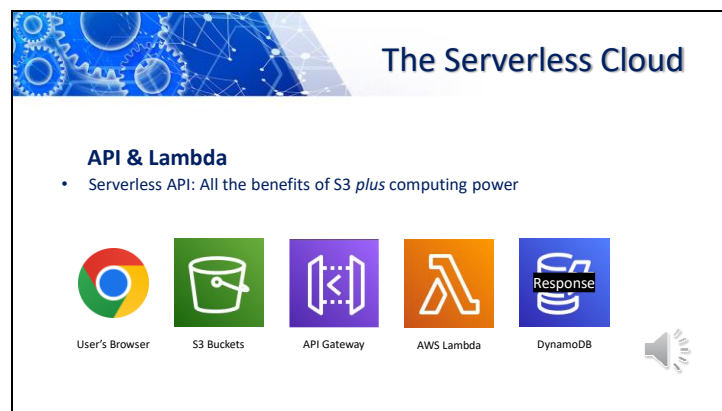
- Serverless API: All the benefits of S3 *plus* computing power

This slide is identical to Slide 13, showing the same serverless architecture flow: User's Browser -> S3 Buckets -> API Gateway -> AWS Lambda -> DynamoDB. A speaker icon is present at the bottom right.

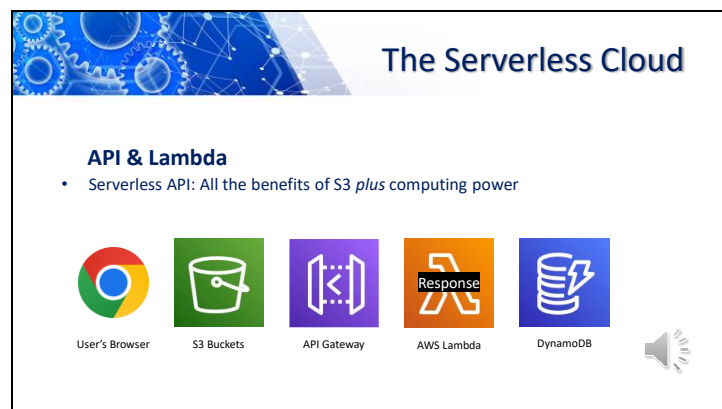
Slide 16



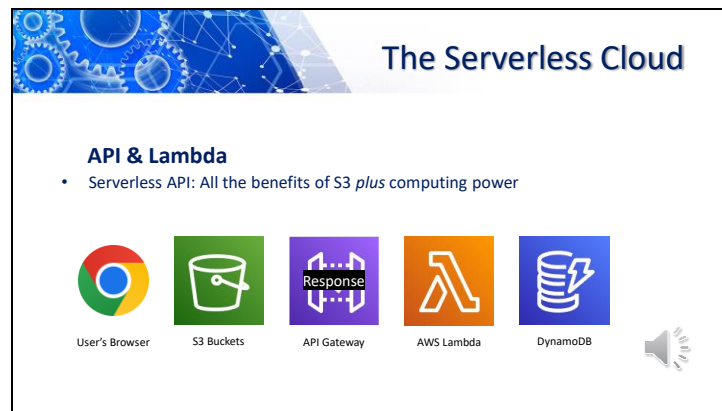
Slide 17



Slide 18



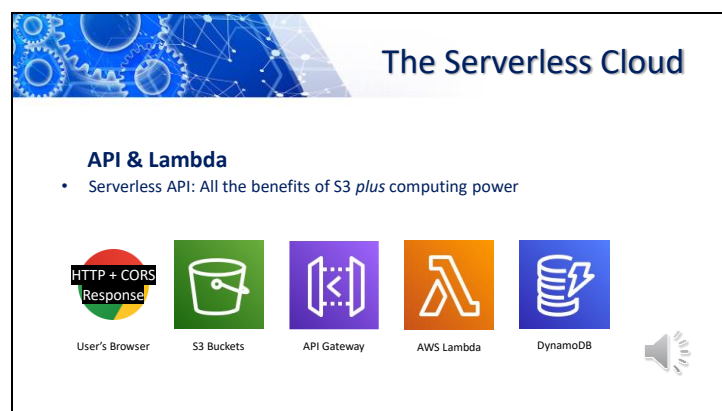
Slide 19



Slide 20



Slide 21



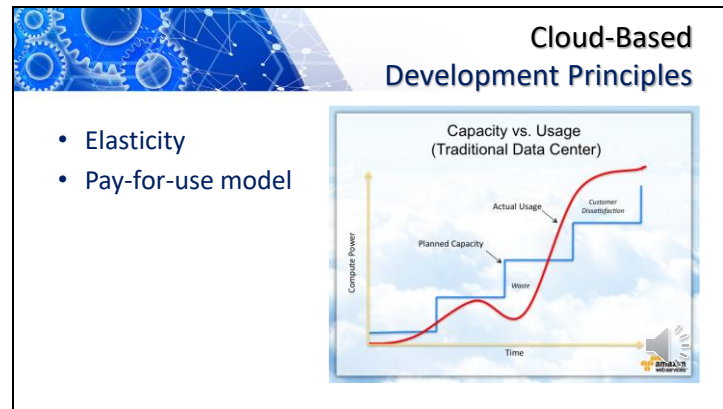
Slide 22



MongoDB has a multi-table document structure, while DynamoDB has a single-table key-value structure. In DynamoDB, the table automatically scales as items are added, columns are created as needed and unique id strings identify every item.

We query the question and answer items for our website using the unique id keys, with answers referencing the question's key to link the two.

The scripts that manage this database manipulation are Lambda functions, such as TableScan for GETting all the items, findOne for finding the specific answers for a single question, Upsert functions for POST and PUT requests, and DeleteRecord for DELETE functionality.



Elasticity describes the need for compute power to adjust in real-time based on user demand. The capacity can be planned ahead of time, and servers can be turned on or off as needs change, but a smaller server may not be able to handle large spikes in user demand. AWS manages huge amounts of server hardware, and can quickly reroute requests to servers with higher capacity to meet demand, then limit that server's production when demand tapers off to reduce waste.


The pay-for-use model is a way of giving the AWS user value, so that a smaller developer only pays for the small amount of power they're using, and a large company with thousands of requests can pay for their usage. There is a cost analysis to be done, but generally the pay-as-you-go is going to provide optimal value for any project that can't afford a server maintenance employee, doesn't have steady demand on their server, or is just starting out and doesn't know what their capacity needs will be.



AWS automatically sets many security features into their products. AWS Web Application Firewall automatically prevents SQL injection, cross-site scripting, and DDoS attacks. The Identity and Access Management, IAM, features allow the developer to set permissions for individuals or groups, allowing for the principle of least privilege to be easily maintained, and supporting multi-factor authentication to prevent common user-side access abuse like phishing or keylogging.


With IAM allowing us to assign roles to users, those roles let us dictate the security policies that allow users to access the data. Roles and policies are complex, and with the LabRole assigned by AWS we were able to do a lot within the lab environment. Unfortunately, some IAM features were not available so my understanding of the policies and roles is limited.

AWS security has many capabilities. For securing API Gateway and Lambda, AWS offers SSL certificates, the web application firewall, MFA, request throttling, and API key options. AWS expects the user to maintain control over resources using these tools, while they protect the infrastructure. AWS automatically protects S3 and database data with Amazon GuardDuty and Amazon Macie, as well as a wide range of security logging capabilities.



Conclusion

Reduce Cost	Scalability	Security
<ul style="list-style-type: none">• No initial investment• No time spent maintaining hardware architecture	<ul style="list-style-type: none">• Automatic request routing• Data redundancy• Access tiers	<ul style="list-style-type: none">• Automatic encryption• Firewall• IAM roles• API keys



Lifting and shifting a MEAN stack web application into AWS microservices was a quick and easy process. Using the serverless architecture that AWS offers is a way to save time and money, particularly for smaller projects and startups. The initial investment is virtually zero, which is a big jump from the time and money invested with traditional server-required projects. Time can be spent developing the product, rather than maintaining server architecture.

AWS scales automatically, routing requests through their expansive server network as necessary while only charging for what is used. As user demand changes, the backend changes with it. Similarly, AWS stores data in multiple places, allotting more space as more data is added. Data is moved through the tiers of storage to make the more volatile data more easily accessible.

AWS security is robust, with all the industry-standard automatic protections such as data encryption and web application firewall, as well as user-implementable protections such as IAM roles and API keys.